



A Tour of GHC's Garbage Collection

Jost Berthold, FP Syd 2024 07
24


TOC

Background: Types of Garbage Collection

GHC's Garbage Collection

Parallel Garbage Collection

Non-Moving GC Extension



**Featuring *heaps ...
of slides
borrowed from the
internet***





Basic algorithms

—The garage metaphor—

Reference counting: Maintain a note on each object in your garage, indicating the current number of references to the object. When an object's reference count goes to zero, throw the object out (it's dead).

Mark-Sweep: Put a note on objects you need (roots). Then recursively put a note on anything needed by a live object. Afterwards, check all objects and throw out objects without notes.

Mark-Compact: Put notes on objects you need (as above). Move anything with a note on it to the back of the garage. Burn everything at the front of the garage (it's all dead).

Copying: Move objects you need to a new garage. Then recursively move anything needed by an object in the new garage. Afterwards, burn down the old garage (any objects in it are dead)!




Background: Types of Garbage Collection (GC)

- Identifies items allocated in memory that are not used any more
- Reclaims allocated memory to store new items
- *Reference counting vs tracing strategies*
 - Reference counting: identifies and removes *dead* items
 - Tracing: identifies and protects items that are *alive*



Terminology

Single threaded collection



Parallel collection



Concurrent collection



Incremental collection



On-the-fly concurrent collection



User program

mutator

collections

and all combinations...



Basic Garbage Collection in GHC

- (Cheney-style) 2-space copying GC
- Heap contains *closures* of evaluated or unevaluated (thunk) data
-

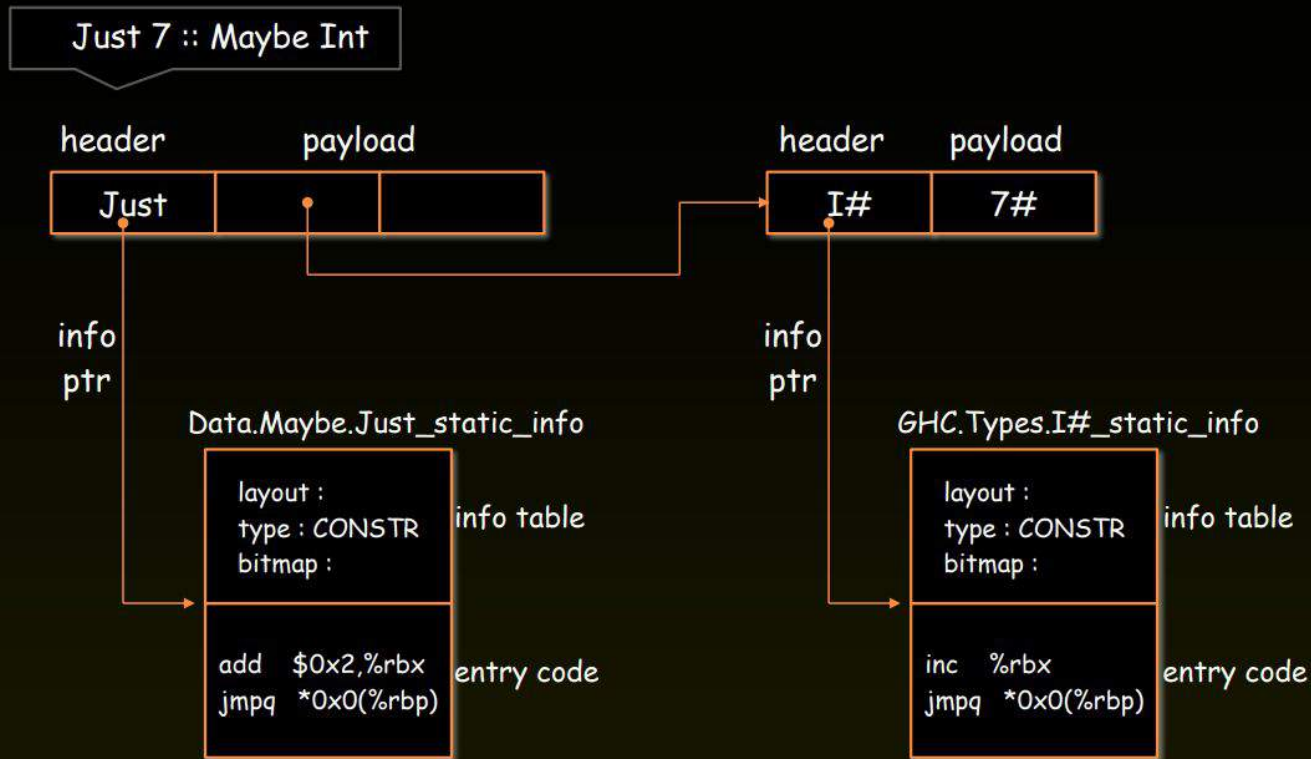
```
./my-program +RTS -H2G -RTS # 2G min heap
```

```
./my-program +RTS -M12G -RTS # 12G max heap
```

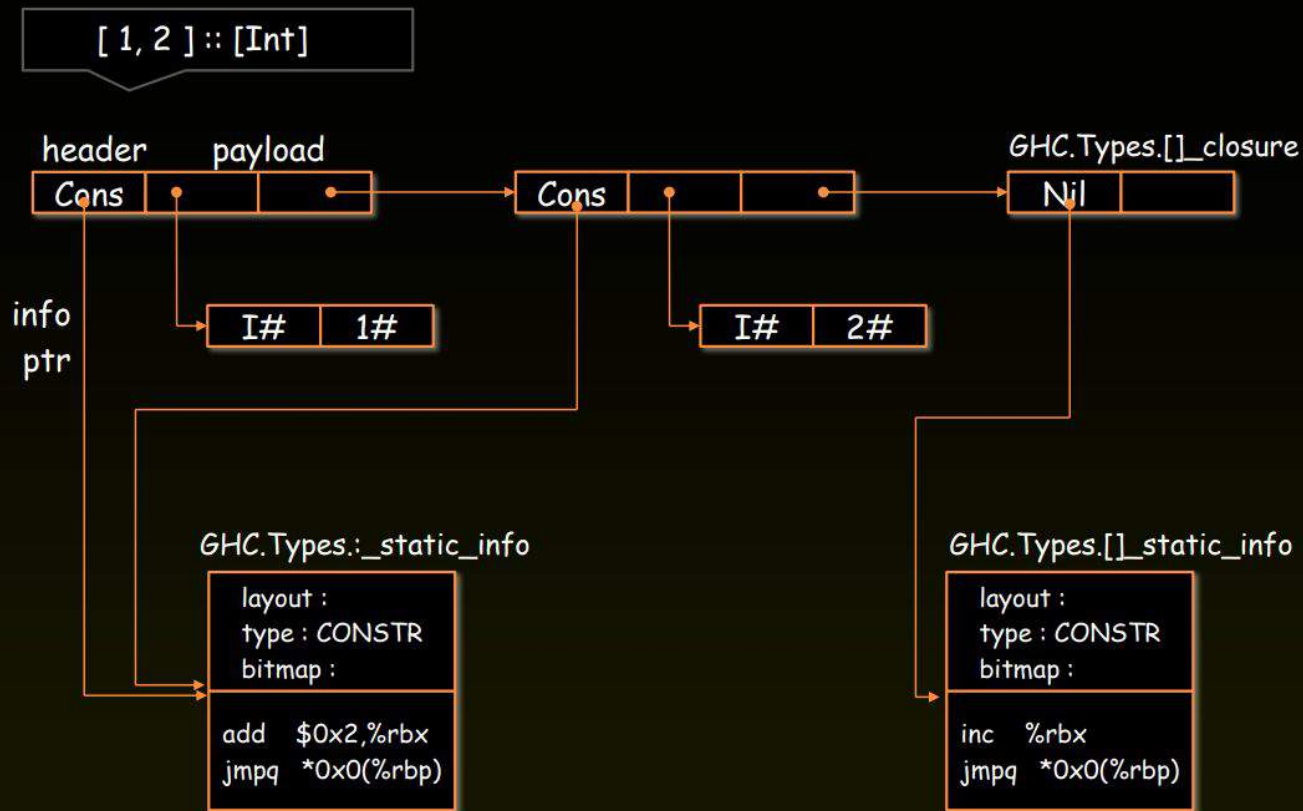
```
./my-program +RTS -m5 -RTS # ensure 5% heap remain available
```

```
./my-program +RTS -sstderr -RTS # get a GC summary at the end
```

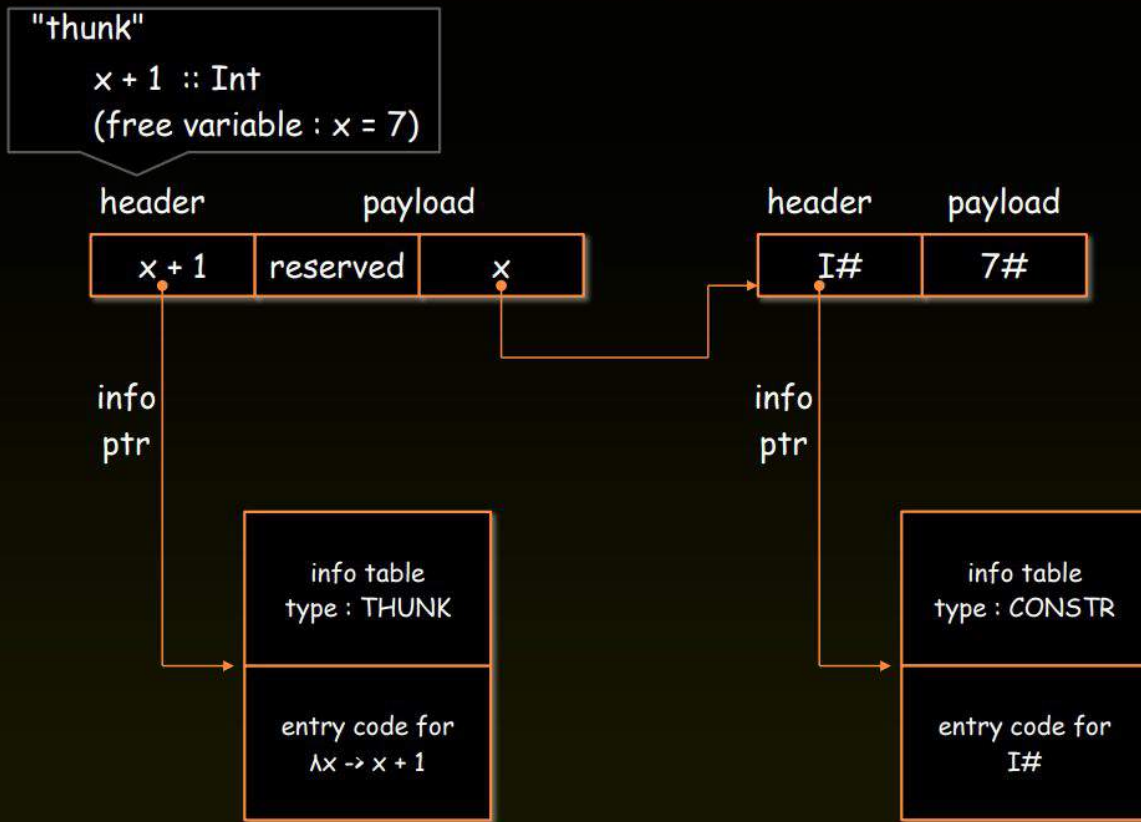
Closure examples : Maybe




Closure examples : List




Closure examples : Thunk




Moving garbage collection




Moving garbage collection




Moving garbage collection: Evacuate roots




Moving garbage collection: Evacuate roots




Moving garbage collection: Evacuate roots



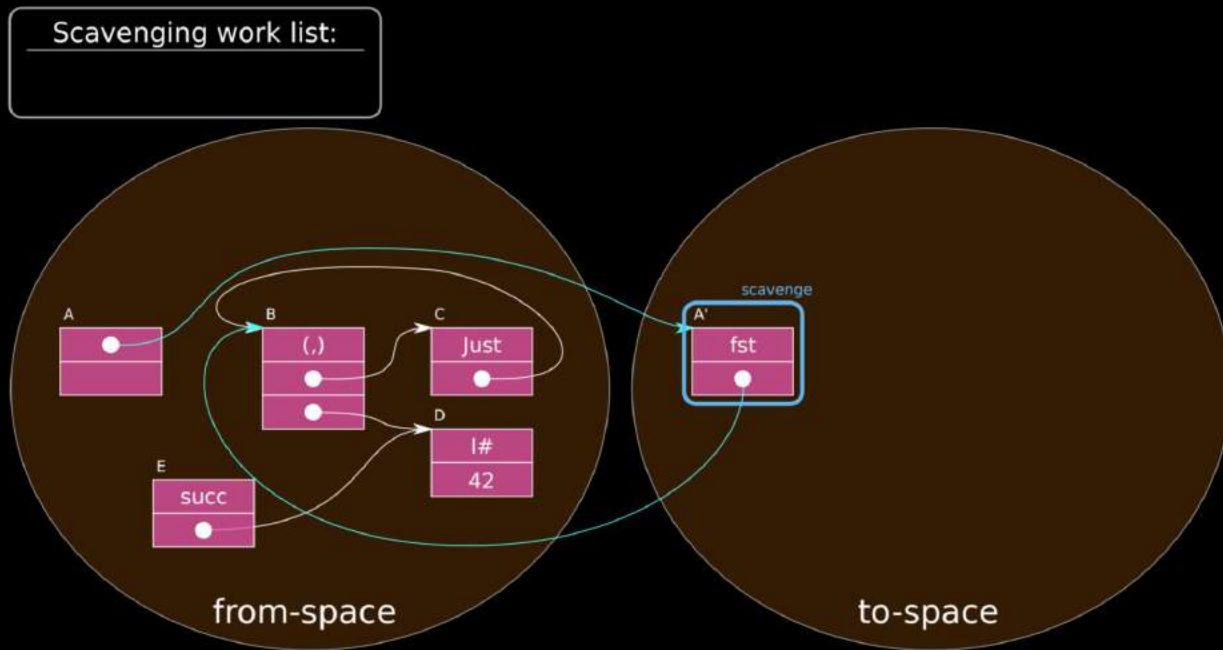
Moving garbage collection: Evacuate roots




Moving garbage collection



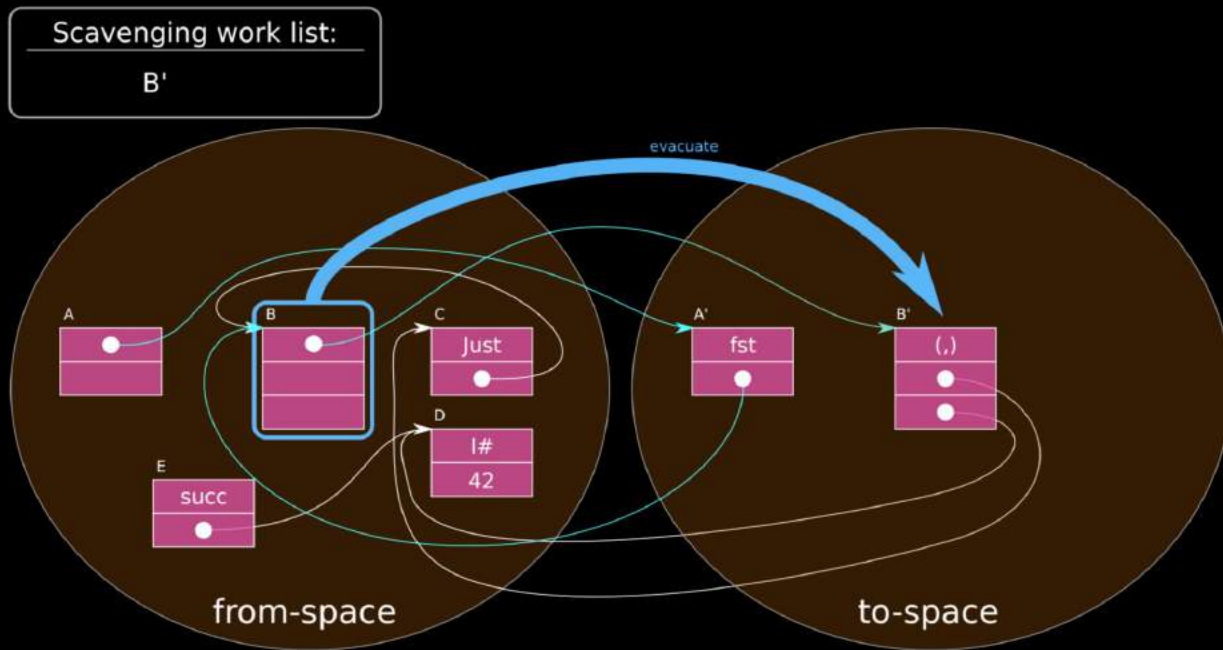
Moving garbage collection: Scavenge A



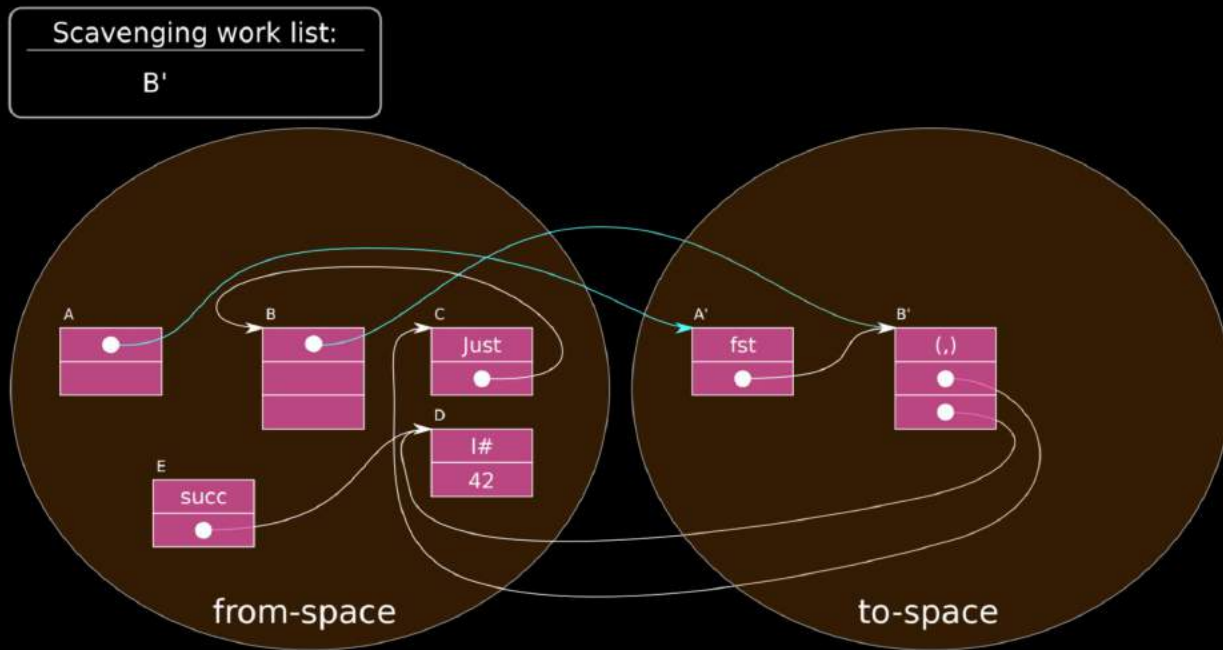
Moving garbage collection: Evacuate B




Moving garbage collection: Evacuate B




Moving garbage collection: Evacuate B




Moving garbage collection




Moving garbage collection: Scavenge B




Moving garbage collection: Evacuate C




Moving garbage collection: Evacuate C




Moving garbage collection: Evacuate C




Moving garbage collection: Scavenging B (cont'd)




Moving garbage collection: Evacuate D




Moving garbage collection: Evacuate D




Moving garbage collection: Scavenging C




Moving garbage collection: Scavenging C



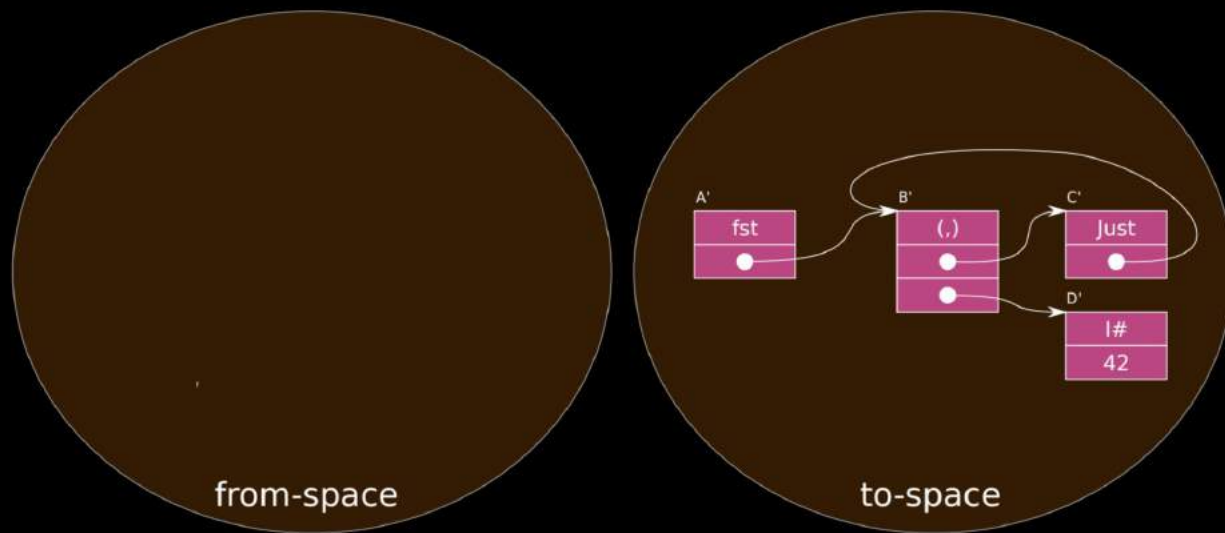
Moving garbage collection: Scavenging D



Moving garbage collection: Finished!



Moving garbage collection: Finished!





Actual (Generational) GHC Garbage Collector

- Several “generations” of items
- Generation 0 is collected often
- Heap objects get *promoted* to higher generations
- Generation 1 is collected when gen.0 GC does not reclaim enough space
- In practice: usually 2 generations, Gen 0 split into 2 steps.



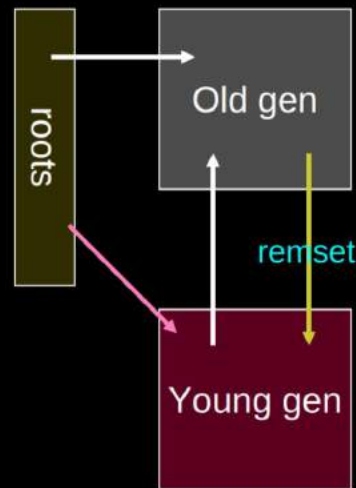
Generational GC

Weak generational hypothesis

- “Most objects die young” [Ungar, 1984]
- Common for 80-95% objects to die before a further MB of allocation.

Strategy:

- Segregate objects *by age* into **generations (regions of the heap)**.
- Collect different generations at different frequencies.
 - so need to “remember” pointers that cross generations.
- Concentrate on the nursery generation to reduce pause times.
 - full heap collection pauses 5-50x longer than nursery collections.





from space

Nursery

Generation 1

to spaces

Slides from Edward Yang 2015




Tenuring

Nursery



from space




Nursery



to spaces

Nursery



Generation 1



Nursery

Generation 1


```
mk_exit()
  entry:
    Hp = Hp + 16;
    if (Hp > HpLim) goto gc;

    v::I64 = I64[R1] + 1;

    I64[Hp - 8] = GHC_Types_I_con_info;
    I64[Hp + 0] = v::I64;

    R1 = Hp;
    Sp = Sp + 8;
    jump (I64[Sp + 0]) ();


gc: HpAlloc = 16;
    jump stg_gc_enter_1 ();
}
```





What about Purity?

→ Write Barriers


→ Parallel Garbage Collection








Minor GC





Why is generational GC
hard? This.

Why is generational GC
hard in Java? This.


Purity to the rescue

- Mutation is rare
- IORefs are slow anyway
- Laziness is a special kind of mutation

Nursery

Generation 1

thunk



Promotion

Nursery

Generation 1

ind

result

immutable now

GC, nursery, generation, aging, promotion

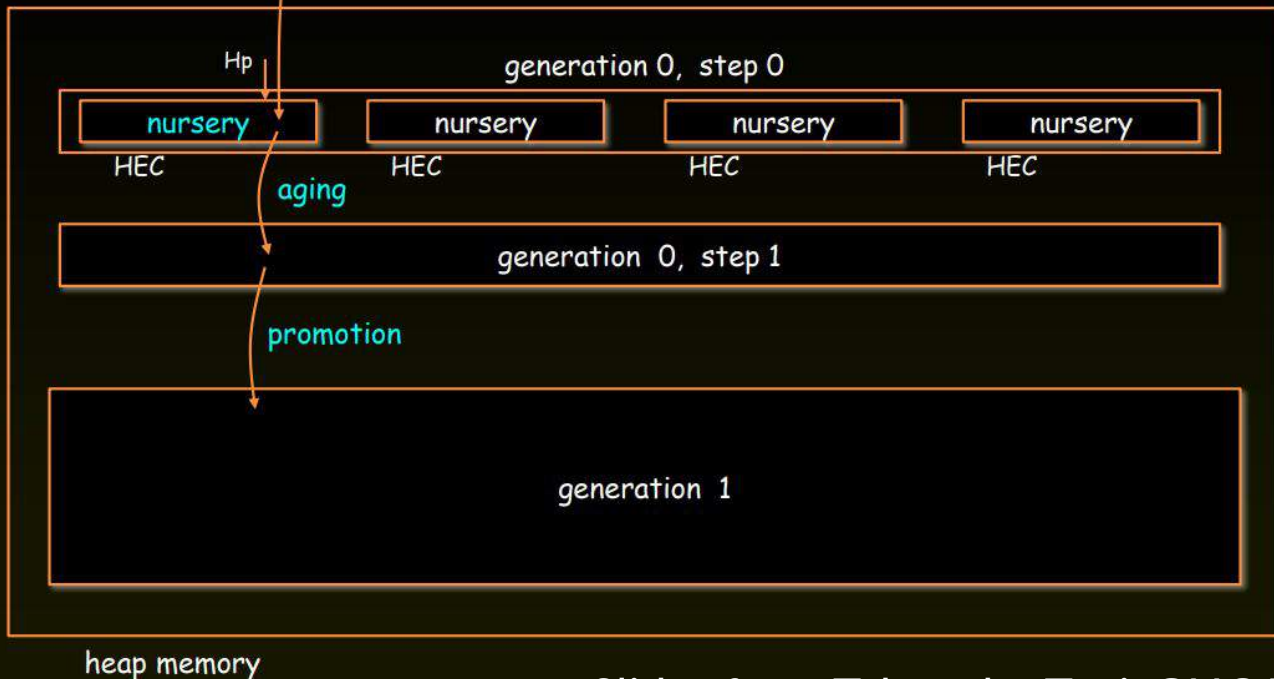
STG land
(Haskell land)

Haskell
Thread

```
./my-program +RTS -G3 -RTS # use 3 generations  
./my-program +RTS -F4 -RTS # live-data factor to control old gen. GC  
./my-program +RTS -A42M -RTS # use 42M for nursery (gen0 step0)
```

Runtime System

allocate







Parallel GC and Compacting GC

- Parallel GC
 - Copying GC in distinct *blocks* of the heap in parallel
 - Can use load-balancing
 - Still stop-the-world (*not concurrent*)
- Compacting GC:
 - Collect the oldest gen. using a mark-sweep-compact collector
 - May reduce memory pressure for programs with large heaps
 - Still stop-the-world (*not concurrent*)

Parallel GC

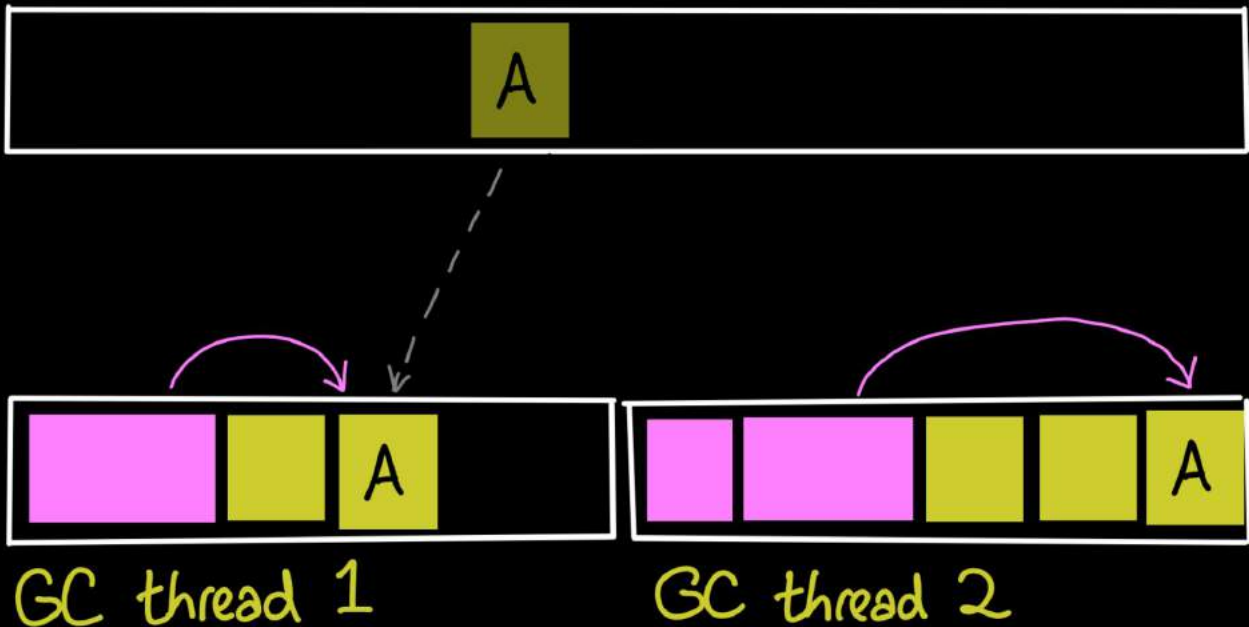
Idea: Split heap into blocks,
and parallelize the scavenging process





Needs synchronization

If A is immutable...

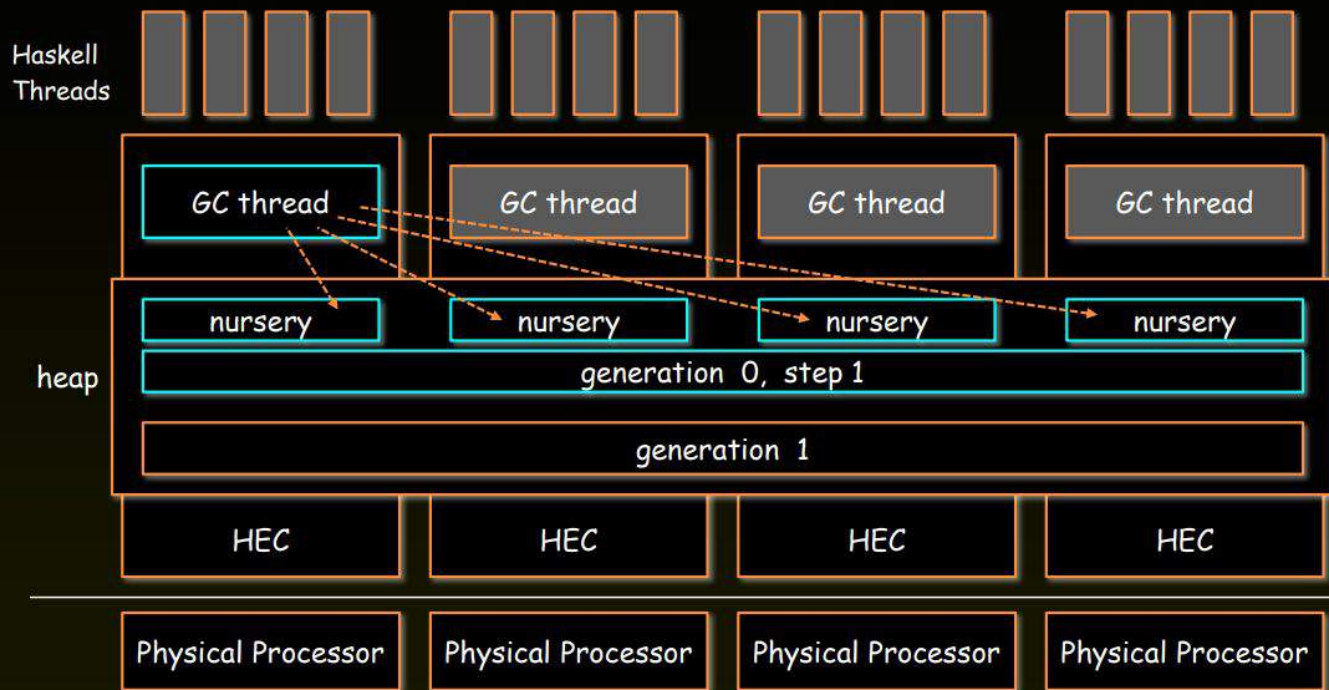


...observationally indistinguishable!

Threads and minor GC

sequential GC for young generation (minor GC)

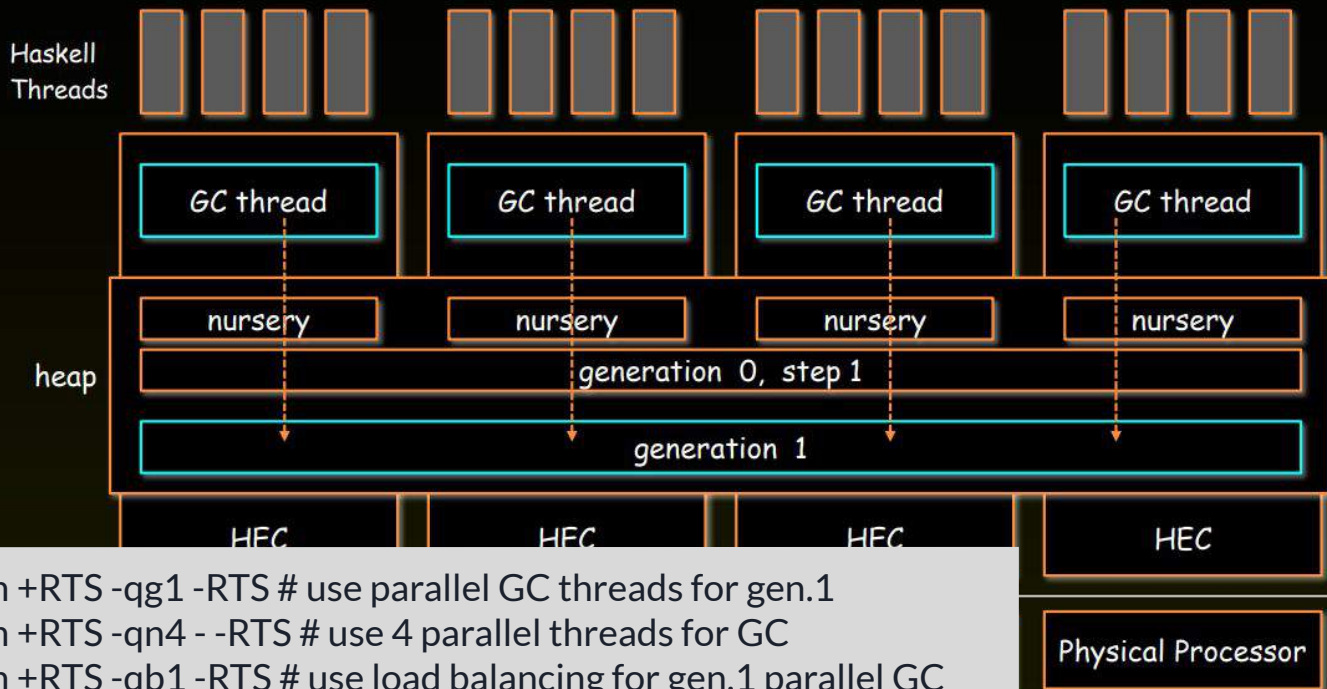
"stop-the-world" GC



Threads and major GC

parallel GC for oldest generation (major GC)

"stop-the-world" GC




```
./my-program +RTS -qg1 -RTS # use parallel GC threads for gen.1
./my-program +RTS -qn4 -RTS # use 4 parallel threads for GC
./my-program +RTS -qb1 -RTS # use load balancing for gen.1 parallel GC
./my-program +RTS -qb -RTS # disable load balancing
```



New: Non-Moving GC


- Advantages of Copying GC:
 - Efficient allocation, data locality, no fragmentation
- Disadvantages
 - Cannot run incrementally nor concurrently
 - Cannot collect a fraction of the heap without scanning all of it
 - Global stop of unknown length (latency)
- Controlling the latency:
 - Incremental collection (bounded amount of work)
 - Concurrent collection (while mutators are running)
- *Here:* only for the old(-est) generation

Hybrid moving/non-moving collector




- ▶ Mutator allocates into a moving nursery
- ▶ Young generation collector evacuates into non-moving heap
- ▶ Non-moving heap collected with mark & sweep
- ▶ Mark & sweep amenable to both incremental and concurrent collection


Mark & Sweep Garbage Collection




Mark & Sweep Garbage Collection: Collect roots




Mark & Sweep Garbage Collection: Marking (A)




Mark & Sweep Garbage Collection: Marking (B)




Mark & Sweep Garbage Collection: Marking (C)




Mark & Sweep Garbage Collection: Marking (D)



Mark & Sweep Garbage Collection: Marking (B)



Mark & Sweep Garbage Collection: Sweep

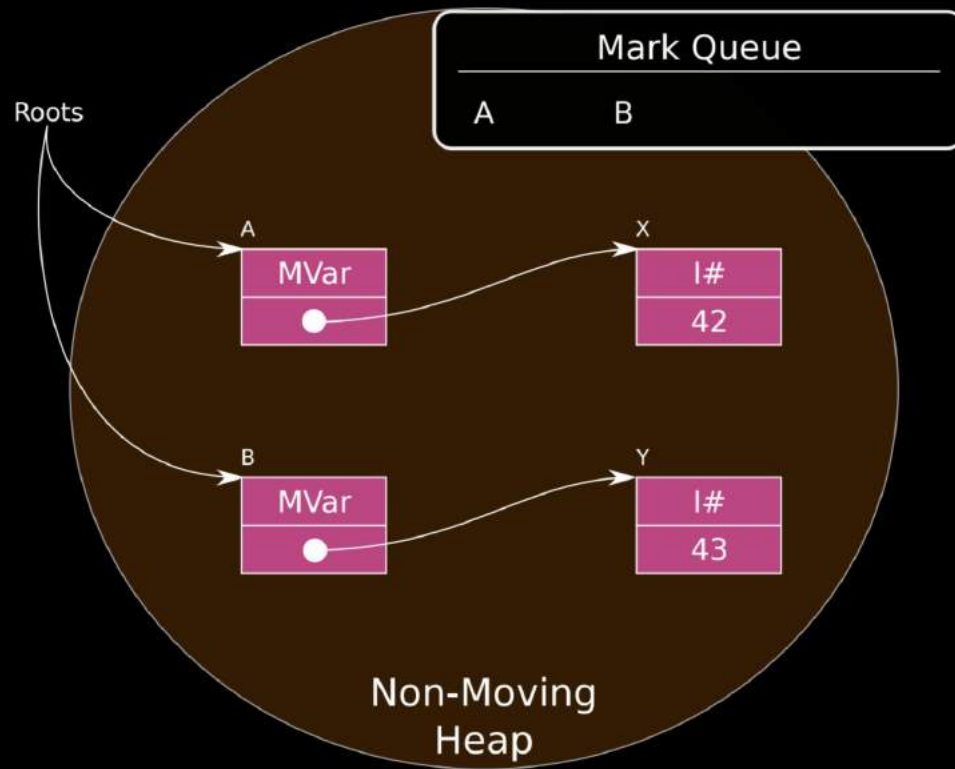




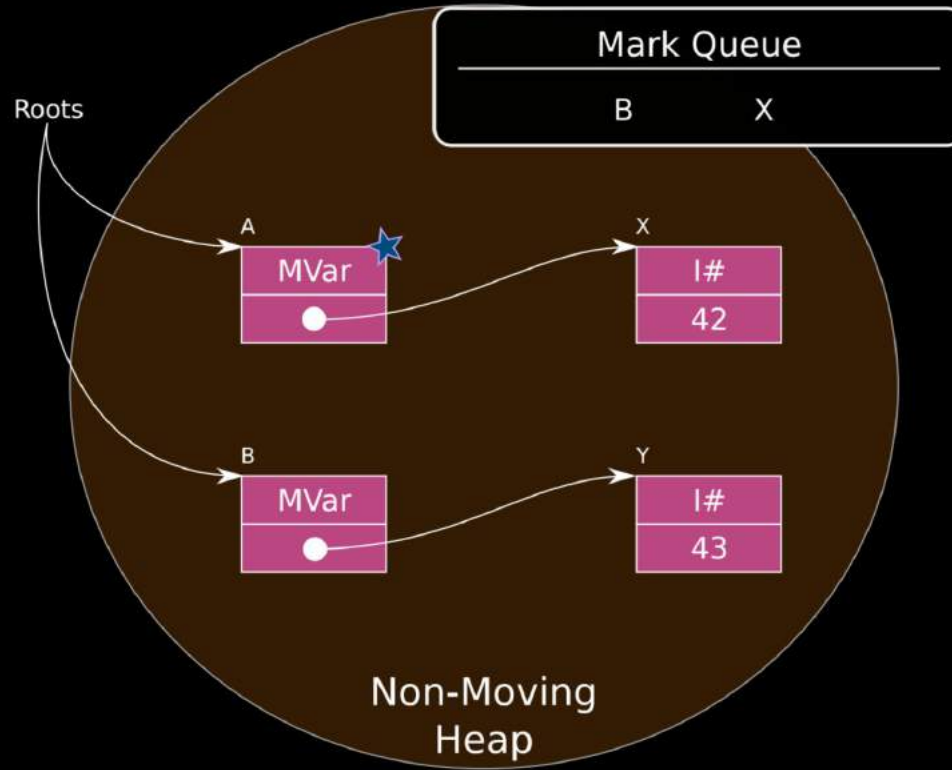
Nonmoving GC for the oldest generation

- References remain valid (nothing is moved/copied)
- Sweep phase leaves a fragmented heap behind
 -
 - Needs more management for the allocator
- Mark/Sweep runs (mostly) concurrently with mutator(s)

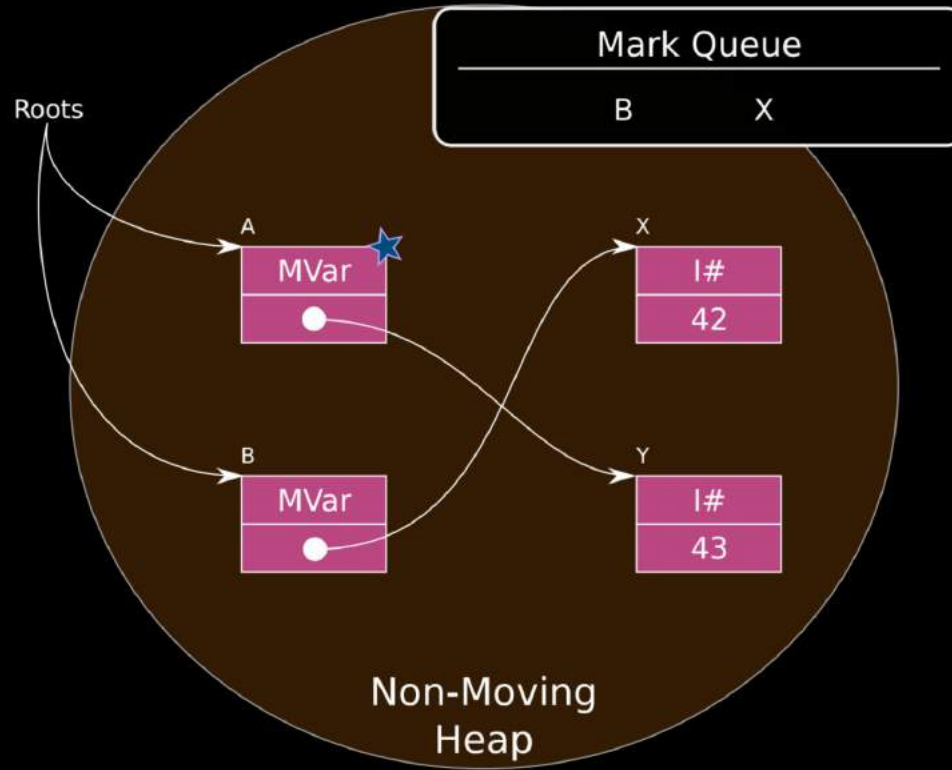
The problem of mutation



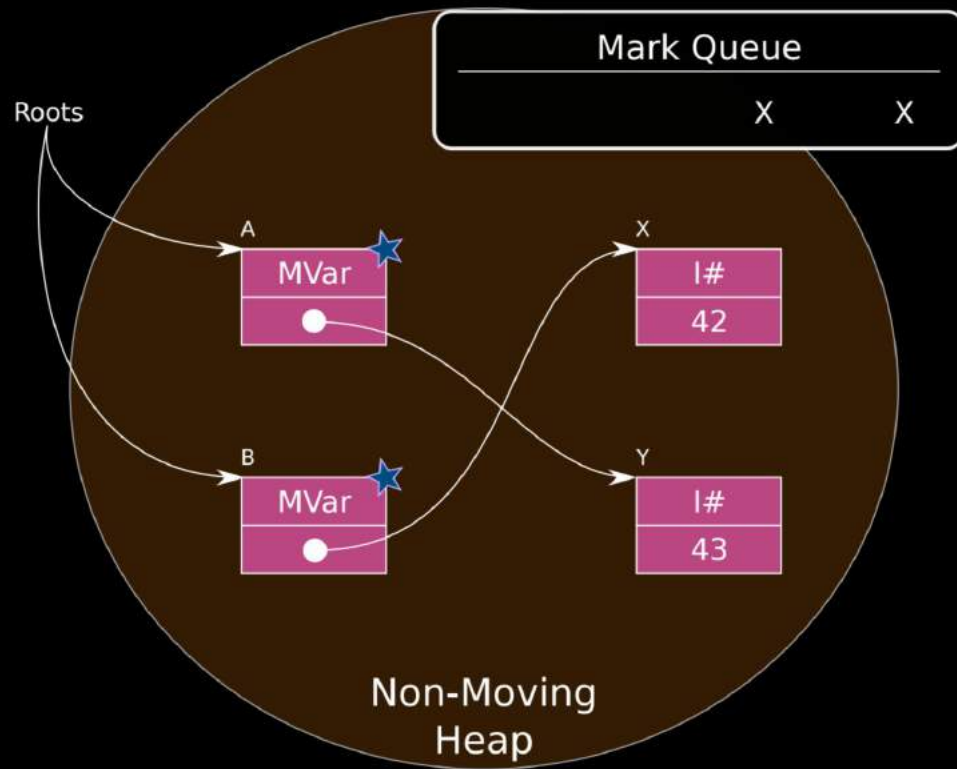
The problem of mutation



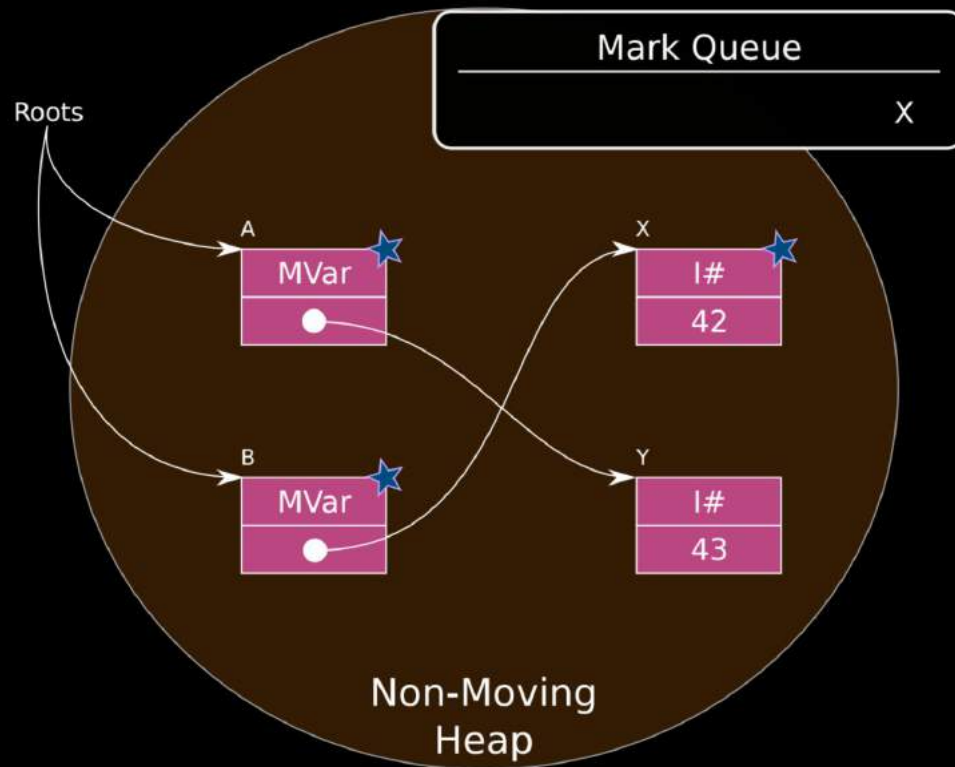
The problem of mutation



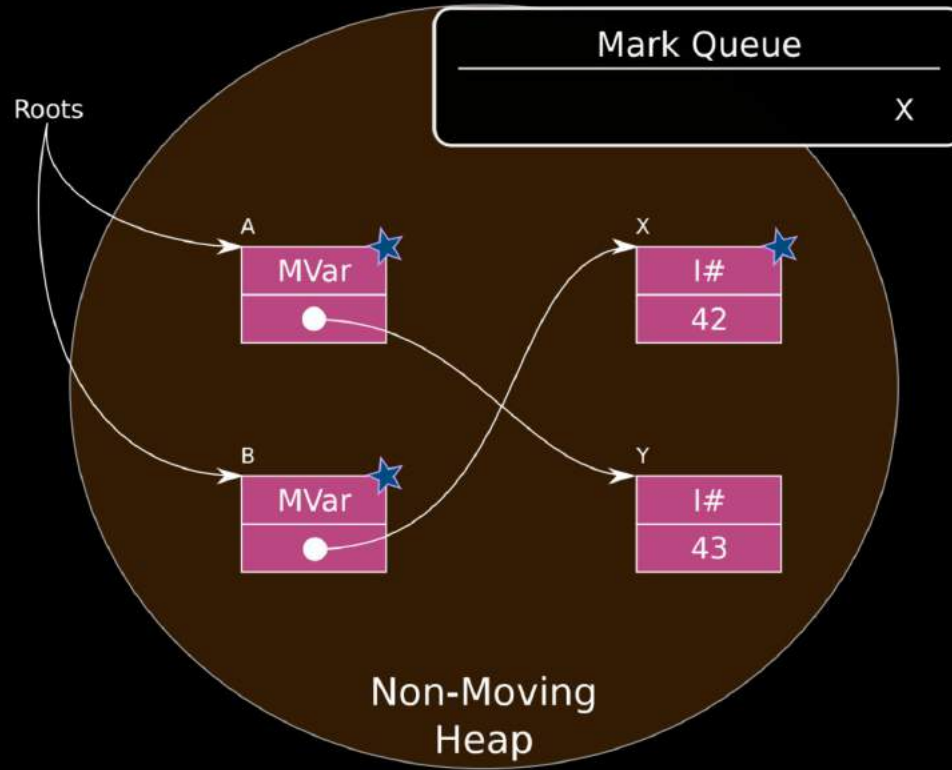
The problem of mutation



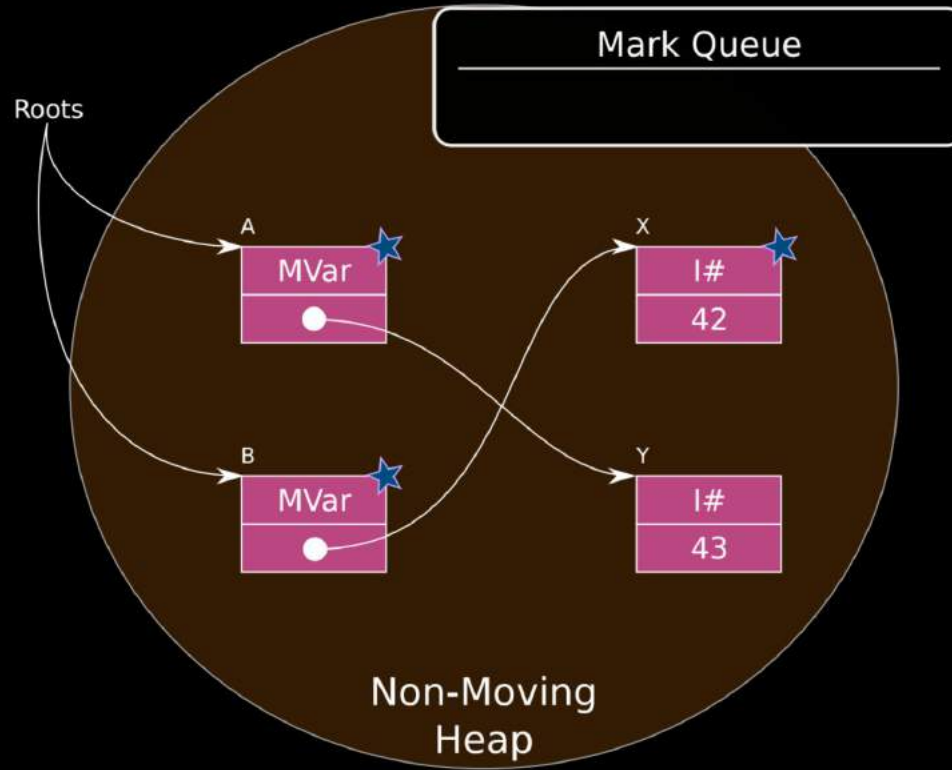
The problem of mutation



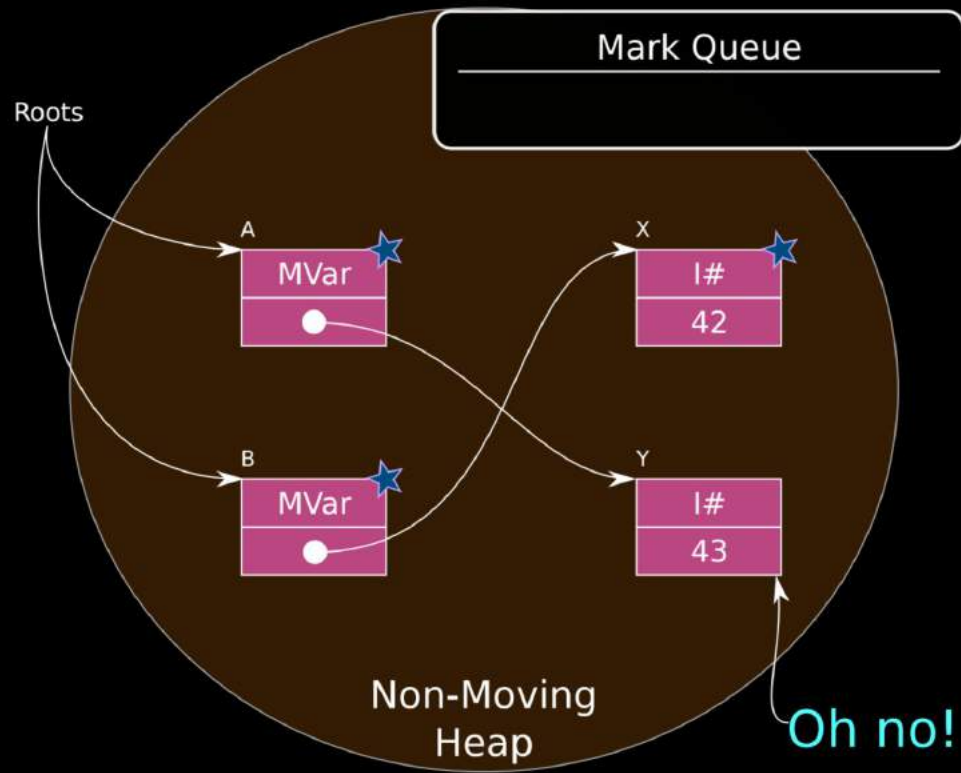
The problem of mutation



The problem of mutation



The problem of mutation



Snapshot-at-the-beginning

Solution:

- ▶ Collect with respect to the state of heap at start of mark (time t_0).

Concretely, the collector must ensure this property (henceforth the *snapshot invariant*):

The collector must mark all objects reachable at t_0 .

N.B. it is also safe to mark objects that were *dead* at t_0 .

Snapshot-at-the-beginning

The snapshot invariant:

The collector must mark all objects reachable at t_0 .

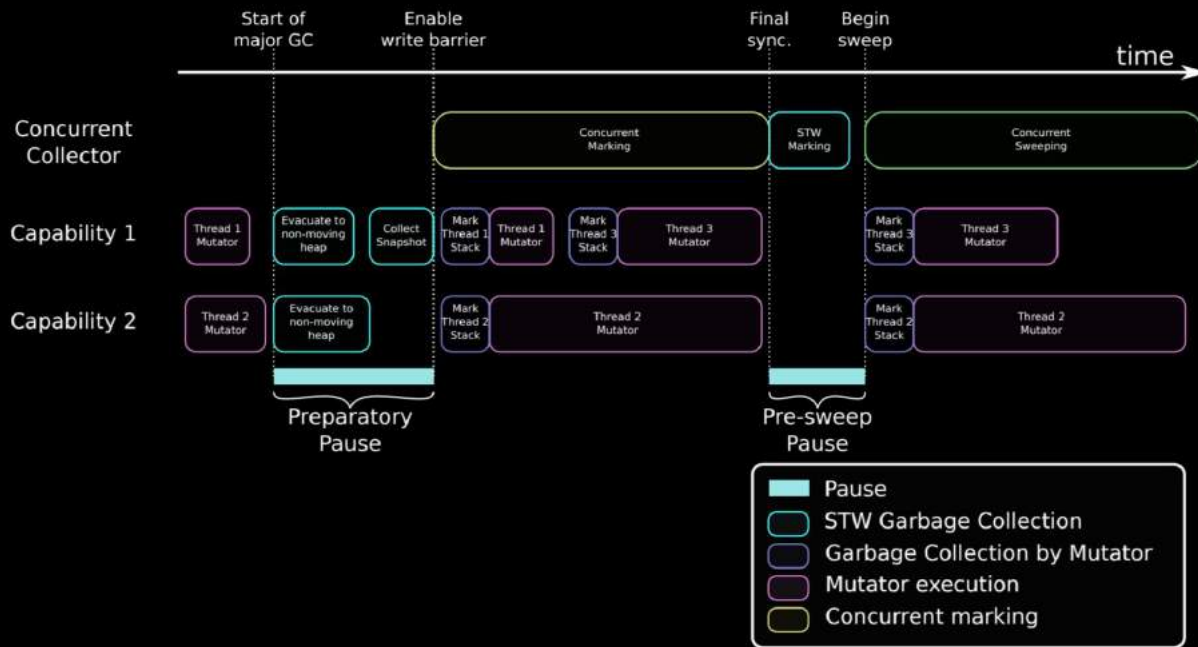
Consequently,

1. All objects live at t_0 will be retained.
2. Many objects dead at t_0 will be freed.
3. All objects allocated after t_0 will be retained.


How to accomplish this?

A write barrier.

The lifecycle of a GC



Allocator structure: The segment



ISMM '20, June 16, 2020, London, UK


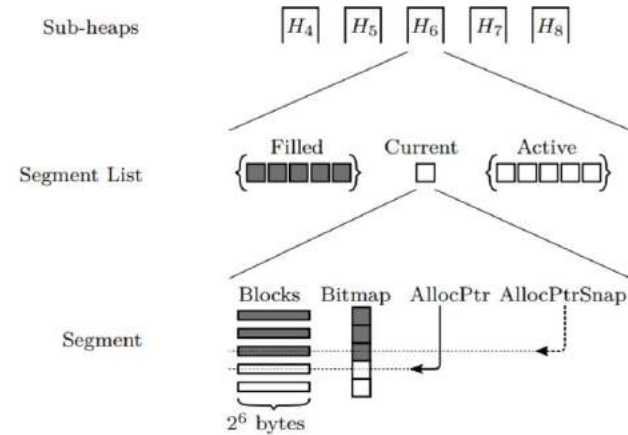



Figure 1. Structure of an allocation heap

From Ben Gamari 2018, Gamari/Dietz2020, UenoOho2015

Determining object allocation state

- ▶ Recall: the snapshot invariant only requires that we mark objects which were reachable at t_0
- ▶ How do we know whether an object was allocated at t_0 ?



- ▶ Record the value of each segment's `next_free` field when the snapshot is taken.
- ▶ During collection we can conclude that objects above the snapshot needn't be marked.

What the new collector won't do...

Concurrent collection isn't a silver bullet:

- ▶ It probably won't make your program run faster
- ▶ It won't make your program scale more effectively across cores (but this may be future work)
- ▶ It may reduce your program's memory footprint, but not by as much as you might expect
- ▶ It is not provide hard realtime latency guarantees
- ▶ It does not mow your lawn (yet)


```
./my-program +RTS -xn -RTS
```

```
./my-program +RTS --nonmoving-gc -RTS
```



Thank you!

Lorem ipsum dolor sit amet,
consectetur adipiscing elit.
Curabitur eleifend a diam
quis suscipit.





References

Slides stolen from:

- [Edward Z Yang](#) (Lecture about GHC RTS)
- [Richard Jones](#) (talk about GC and multicore)
- [Ben Gamari](#) (-early- design of the nonmoving GC)
- [Takenobu Tani](#) (GHC RTS illustrated)

Good Reads:

- [Short paper](#) / [Paper nonmoving GC](#)
- [Prior Art](#) (SML#, not Haskell)
- [A summary of modern GC research](#)
- (Original) [Generational GC for GHC](#)
- [Incremental GC Experiment in GHC](#) and [successive further work](#)
- [Parallel GC](#)
- [Options in GHC manual](#)